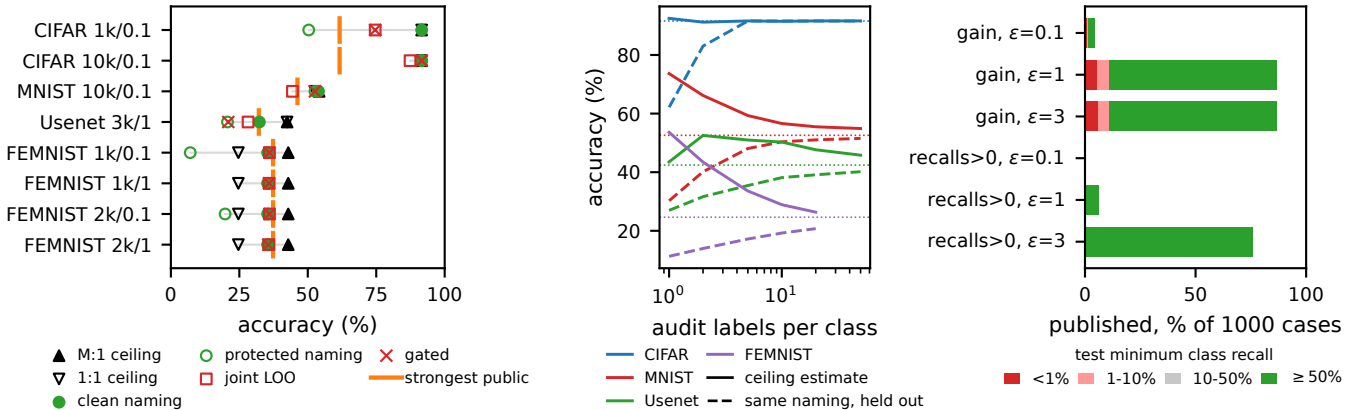
(a) loss decomposition
CIFAR 1k/0.1
CIFAR 10k/0.1
MNIST 10k/0.1
Usenet 3k/1
FEMNIST 1k/0.1
FEMNIST 1k/1
FEMNIST 2k/0.1
FEMNIST 2k/1
0
25
50
75
100
accuracy (%)
M:1 ceiling
1:1 ceiling
clean naming
protected naming
joint LOO
gated
strongest public
(b) auditing the ceiling
accuracy (%)
80
60
40
20
10^0
10^1
audit labels per class
CIFAR
MNIST
Usenet
FEMNIST
ceiling estimate
same naming, held out
(c) what the rules publish
gain, ε=0.1
gain, ε=1
gain, ε=3
recalls>0, ε=0.1
recalls>0, ε=1
recalls>0, ε=3
0
50
100
published, % of 1000 cases
test minimum class recall
<1%
1-10%
10-50%
≥ 50%